


CAISSA-JEPA | Complete legal-reply development study


Development screen not passed | Prespecified candidate: JEPA tail

One global rate per family selected on EXACT regret; unchanged for HYBRID. Full labels, three paired training seeds.


A EXACT: equal-game mean




B EXACT: per game




C Raw-tail effects; saved 95% intervals




D HYBRID: equal-game mean



E HYBRID: per game



F Raw-tail effects; saved 95% intervals





Adaptive DEVELOPMENT on 209 reused roots (107 Connect4, 102 Reversi), with seeds 17/29/43. Bars are three-seed means; ticks show those seeds. Paired root-bootstrap intervals use 10,000 draws, shared across the fixed seeds and stratified by game. Intervals do not correct adaptive model selection and are not confirmation. Same legal tree does not imply equal wall time or active compute. All 42 cells and both rates appear on the appendix page; no model scoring is performed by this plotter.

CAISSA-JEPA | Complete 42-cell development inventory

Every family x both rates x seeds 17/29/43; no cell omitted. Each entry is saved per-cell action regret.

Asterisk marks the EXACT-selected rate used unchanged in both tracks. Candidate rows are explicitly labeled.



Adaptive development, not confirmatory evidence. The fixed color scale spans the full zero-sum action-regret range. Equal-game cells average the two saved game means, so unequal game counts do not change their weight. No confidence interval is inferred for an individual cell. All cells are complete and uncollapsed under the strict report.